

PRISM-EM for Non-Stationary Poisson GLMs from Spike Trains: Model, Objective, and Implementation. Ver3

Jan Balewski, NERSC

March 23, 2026

Abstract

This document describes the current `prism_EM_train.py` implementation for fitting a non-stationary Poisson GLM with latent state mixing (PRISM-EM) ¹. The model uses a shared connectivity matrix across states, state-specific bias vectors, and time-varying simplex coefficients. Training is performed by block coordinate descent: a sequential projected-gradient E-step for coefficients followed by Viterbi decoding to one-hot state assignments, and an Adam-based M-step for model parameters with off-diagonal L1 shrinkage and spectral-radius projection. Spectral-radius enforcement, L1 pruning, and learning-rate decay are each activated after a configurable number of EM iterations. We summarize the exact objective used in code, the decoding method, all current default hyperparameters, and initialization modes (data-driven vs random) for A , B , and C .

Contents

1	EM Training for Non-Stationary Poisson GLM	2
1.1	Introduction and Model Assumptions	2
1.2	Objective Function	2
1.3	Why Block Coordinate Descent (EM) over Joint Optimization	2
1.4	Algorithm: Block Coordinate Descent (EM)	3
1.4.1	E-step: Coefficient Inference via Simplex PGD, then One-Hot Encoding	3
1.4.2	M-step: Dictionary Update via Adam	3
1.5	State Decoding	4
2	Default Hyperparameters and Runtime Settings	4
3	Initialization of Model	4
3.1	State-Coefficient Initialization ($C \equiv \{c_t\}$)	4
3.2	Bias Initialization (B)	5
3.3	Connectivity Initialization (A)	6

¹PRISM: PRobabilistic Inference of Sparse Mixed-state dynamics; EM: Expectations Maximization

1 EM Training for Non-Stationary Poisson GLM

1.1 Introduction and Model Assumptions

We consider a population of N neurons observed over T discrete time bins of width Δt (seconds). The population activity is governed by M latent dynamical states that switch over time. The key structural assumption is that all states share a single connectivity matrix $A \in \mathbb{R}^{N \times N}$ but differ in baseline excitability through per-state bias vectors $B_m \in \mathbb{R}^N$, $m = 1, \dots, M$.

At each time bin t , a coefficient vector $c_t \in \mathbb{R}^M$, constrained to the probability simplex ($c_{mt} \geq 0$, $\sum_m c_{mt} = 1$), selects the effective bias as a convex combination:

$$B_t = \sum_{m=1}^M c_{mt} B_m = c_t^\top \mathbf{B}, \quad (1)$$

where $\mathbf{B} \in \mathbb{R}^{M \times N}$ stacks the M bias vectors row-wise.

Spike counts follow

$$\eta_t = A Y_{t-1} + B_t, \quad \tilde{\eta}_t = \min(\eta_t, \eta_{\text{clip}}), \quad \mu_t = \exp(\tilde{\eta}_t) \Delta t, \quad Y_t \sim \text{Poisson}(\mu_t). \quad (2)$$

Notation summary. N : number of neurons; M : number of latent states; T : number of time bins; Δt : time-bin width (from metadata); η_{clip} : clipping threshold (from metadata); $Y_t \in \mathbb{N}_0^N$: observed spikes; $\mu_t \in \mathbb{R}_+^N$: expected spike count per bin.

1.2 Objective Function

The training objective uses uniformly weighted (unweighted) Poisson NLL in both the E-step and M-step:

$$\mathcal{J} = \underbrace{\sum_{t=2}^T \sum_{i=1}^N [\mu_{t,i} - Y_{t,i} (\tilde{\eta}_{t,i} + \log \Delta t)]}_{\text{Poisson NLL}} + \underbrace{\lambda_2 \sum_{t=2}^T \|\vec{c}_t - \vec{c}_{t-1}\|^2}_{\text{temporal smoothness (E-step only)}} + \underbrace{\lambda_3 \overline{|A_{\text{off}}|}}_{\text{L1 sparsity on off-diagonal (M-step only)}}, \quad (3)$$

with simplex constraints $c_t \in \Delta^{M-1}$ and spectral constraint $\rho(A) \leq \rho_{\text{max}}$, where the off-diagonal L1 term is

$$\overline{|A_{\text{off}}|} = \frac{1}{N(N-1)} \sum_{\substack{i,j=1 \\ i \neq j}}^N |A_{ij}|. \quad (4)$$

In the λ_2 term, the sum over states m is already implicit in the L_2 norm.

1.3 Why Block Coordinate Descent (EM) over Joint Optimization

In principle one could jointly optimize A , $\mathbf{B}$, and $c_{1:T}$ by simultaneous gradient descent on the full objective $\mathcal{J}$. In practice, the EM-style alternation is preferable for three structural reasons.

Breaking the bilinear coupling. The coefficients c_t and the bias matrix $\mathbf{B}$ enter the log-rate linearly through the product $c_t^\top \mathbf{B}$, which then passes through an exponential. When both are free to move, the landscape contains saddle points and flat valleys: a large c_{tm} paired with a small B_m produces the same likelihood as the reverse. Fixing one side at each step eliminates this ambiguity. The Viterbi hard-assignment variant sharpens the separation further: with one-hot $c_t = e_{\hat{s}_t}$, each time bin selects exactly one bias row, so the M-step reduces to M independent Poisson GLM regressions and the bilinear degeneracy vanishes entirely.

Exploiting temporal structure in the E-step. The coefficient vectors $c_{1:T}$ number $O(TM)$ and vastly outnumber the $O(N^2 + MN)$ model parameters. Joint optimization would need to balance learning rates across these disparate groups while enforcing a simplex constraint at every time bin. By contrast, the E-step decomposes into a single forward sweep of cheap M -dimensional simplex-projected gradient updates, naturally incorporating the temporal smoothness penalty $\lambda_2 \|c_t - c_{t-1}\|^2$ through the sequential dependence on the previous accepted c_{t-1} .

Avoiding the log-sum-exp mismatch. Under soft (probabilistic) state assignments the M-step must fit parameters to the mixture rate $\sum_m c_{tm} \exp(\eta_{tm})$, whereas the model parameterizes the log-rate as $\sum_m c_{tm} \eta_{tm}$ inside a single exponential. These two quantities differ whenever the assignment is not one-hot, creating a systematic bias in the gradient signal for $\mathbf{B}$. The hard-EM strategy—Viterbi decode followed by one-hot encoding—eliminates this mismatch by construction, at the cost of losing the monotonic likelihood improvement guarantee of soft EM, a trade-off that is empirically favorable.

1.4 Algorithm: Block Coordinate Descent (EM)

The non-convex problem is solved by alternating:

- E-step: update $c_{1:T}$ with $A, \mathbf{B}$ fixed, then Viterbi-decode to one-hot state assignments.
- M-step: update $A, \mathbf{B}$ with one-hot $c_{1:T}$ fixed.

for K_{EM} outer iterations.

1.4.1 E-step: Coefficient Inference via Simplex PGD, then One-Hot Encoding

Holding $(A, \mathbf{B})$ fixed, coefficients are inferred by one forward sweep from $t = 2$ to T . At each time bin:

1. Build $Z_t[m, :] = A Y_{t-1} + B_m \in \mathbb{R}^{M \times N}$.
2. Run `pgd_iter` projected-gradient iterations:

$$\begin{aligned} \eta_t &= c_t^\top Z_t, \quad \tilde{\eta}_t = \min(\eta_t, \eta_{\text{clip}}), \quad \mu_t = \exp(\tilde{\eta}_t) \Delta t, \\ g_t &= Z_t (\mu_t - Y_t) + 2\lambda_2 (c_t - c_{t-1}), \quad c_t \leftarrow \Pi_{\Delta^{M-1}}(c_t - \alpha_E g_t). \end{aligned}$$

3. Accept c_t , continue to next time bin.

Projection uses the sort-based simplex projection (Duchi et al., $O(M \log M)$).

After the forward sweep completes, the soft coefficients $c_{1:T}$ are decoded into a hard state sequence via Viterbi (Section 1.5). The M-step then receives one-hot encoded state assignments rather than the soft simplex probabilities:

$$\hat{S}_{1:T} = \text{Viterbi}(c_{1:T}), \quad c_t^{(\text{M-step})} = e_{\hat{S}_t}, \quad (5)$$

where e_m is the m -th standard basis vector. This ensures the M-step fits each state’s parameters from cleanly separated data rather than from soft mixtures.

1.4.2 M-step: Dictionary Update via Adam

Holding c_t fixed (as one-hot vectors), update $A, \mathbf{B}$ for K_M epochs using Adam over shuffled mini-batches of (Y_{t-1}, Y_t, c_t) .

Per mini-batch:

$$\ell_M = \frac{1}{|\mathcal{B}|} \sum_{t \in \mathcal{B}} \sum_{i=1}^N [-Y_{t,i} \log(\mu_{t,i} + \varepsilon) + \mu_{t,i}] + \lambda_3 \overline{A_{\text{off}}}, \quad (6)$$

with $\varepsilon = 10^{-8}$.

Then, subject to delay scheduling (operations are activated only after a configurable number of EM iterations have passed):

1. Off-diagonal soft-thresholding (activated after K_{prune} EM iterations):

$$A_{ij} \leftarrow \text{sign}(A_{ij}) (|A_{ij}| - \alpha_M \lambda_3)_+, \quad i \neq j.$$

2. Periodic spectral-radius projection (activated after K_ρ EM iterations, applied every R batches). On each application, in the multi-GPU setting, A and B are first averaged across ranks, then the spectral constraint is enforced, and the result is broadcast:

$$A \leftarrow A \cdot \min\left(1, \frac{\rho_{\max}}{\rho(A)}\right).$$

Learning rate follows linear decay from $\alpha_M^{(0)}$ to $\alpha_M^{(0)} f_{\text{end}}$. The decay is activated only after K_{lr} EM iterations have passed; the total number of decay steps spans the remaining M-step epochs.

1.5 State Decoding

Viterbi decoding is applied both within the EM loop (after each E-step, to produce one-hot assignments for the M-step) and after the final EM iteration (to produce the output state sequence). The decoder uses a homogeneous transition prior:

$$p_{\text{stay}} = \exp\left(-\frac{\Delta t}{\tau_{\text{dwell}}}\right), \quad p_{\text{switch}} = \frac{1 - p_{\text{stay}}}{M - 1}. \quad (7)$$

The decoded sequence is

$$\hat{S}_{1:T} = \arg \max_{\{s_t\}} \sum_t [\log c_{t,s_t} + \log p(s_t | s_{t-1})]. \quad (8)$$

Per-bin confidence is

$$\text{CL}_t = 1 - \max_m c_{t,m}.$$

2 Default Hyperparameters and Runtime Settings

Required input. `--num_states` is required (no default).

Dataset-provided values. Δt (`time_step_sec`) and η_{clip} (`poisson_eta_clip`) are loaded from spike-file metadata.

3 Initialization of Model

Initialization is controlled by:

$$\text{--init_states}, \quad \text{--init_B}, \quad \text{--init_A} \in \{\text{data}, \text{rand}\},$$

with default `data` for all three.

3.1 State-Coefficient Initialization ($C \equiv \{c_t\}$)

random mode:

$$c_t = \frac{1}{M} \mathbf{1}, \quad S_t = -1 \quad \forall t.$$

So coefficients start uniform over states.

Group	Parameter (CLI)	Default
EM structure	<code>--num_em_iters</code>	20
	<code>--m_epochs</code>	5
E-step	<code>--pgd_iter</code>	5
	<code>--lr_estep</code>	0.03
	<code>--lambda2</code>	2.0
M-step	<code>--lr_mstep</code>	3×10^{-3}
	<code>--lr_end_factor</code>	0.1
	<code>--lambda3</code>	0.02
	<code>--rho_max</code>	0.95
	<code>--prescale_m_step_4_ArhoMax</code>	50 (batches)
	<code>--delay_em_iter_4_ArhoMax</code>	3
	<code>--delay_em_iter_4_lrDecay</code>	1
	<code>--delay_em_iter_4_Aprune</code>	1
	<code>--batch_size</code>	2048
	<code>--minW</code>	0.01 (reporting only)
Data window	<code>--time_range_sec</code>	[0.0, 60.0] s
Decoding	<code>--decode_dwell_sec</code>	0.3 s
Init	<code>--init_A</code>	data (or rand)
	<code>--init_B</code>	data (or rand)
	<code>--init_states</code>	data (or rand)

Table 1: Current defaults in `prism_EM.train.py`.

data mode:

1. Aggregate spikes into coarse blocks of length $\lceil \tau_{\text{dwell}} / \Delta t \rceil$ bins.
2. For each block, compute a scalar mean population firing rate.
3. Find $M - 1$ thresholds minimizing within-cluster variance (dynamic-programming segmentation on unique rate values).
4. Assign each block to a state by thresholds.
5. Convert block states to soft probabilities: dominant state gets 0.8, remainder 0.2 split across others, with transition smoothing (2/3 vs 1/3 at boundaries).
6. Map block-level c and S back to original bin resolution.

3.2 Bias Initialization (B)

random mode: No external seed is applied; model keeps random Gaussian initialization from `torch.randn(...)*0`

data mode: For each neuron n :

$$r_n = \frac{1}{T} \sum_t Y_{t,n} / \Delta t, \quad b_n = \log(\max(r_n, 10^{-6})).$$

The same b -vector is assigned to every state row of B .

3.3 Connectivity Initialization (A)

random mode: No external seed is applied; model keeps random Gaussian initialization from `torch.randn(...)*0`

data mode (covariance/OLS): Using first $T_{\max} = 50,000$ bins (or fewer if shorter):

$$A_{\text{ols}} = \left(\sum_t Y_t Y_{t-1}^\top \right) \left(\sum_t Y_{t-1} Y_{t-1}^\top \right)^\dagger.$$

Then all elements of A_{ols} are scaled by a factor of 3:

$$A \leftarrow 3 A_{\text{ols}}.$$

No spectral-radius rescaling is applied at initialization; the spectral constraint is enforced later during training via the delayed projection mechanism.

Diagnostics (condition number, spectral radius, one-step R^2 , Frobenius norm, bins used) are computed on the pre-scaled OLS estimate and saved in metadata.